



Vagrant & Docker

AMT – Labo 2 (more or less)

Authors: Bignens Julien & Brito Carvalho Bruno

Summary

- Introduction
- Architecture
- What we actually do
- Sharing the volumes
- Some magic
- Basic example
- Vagrantfile
- Dockers
- A few tricks



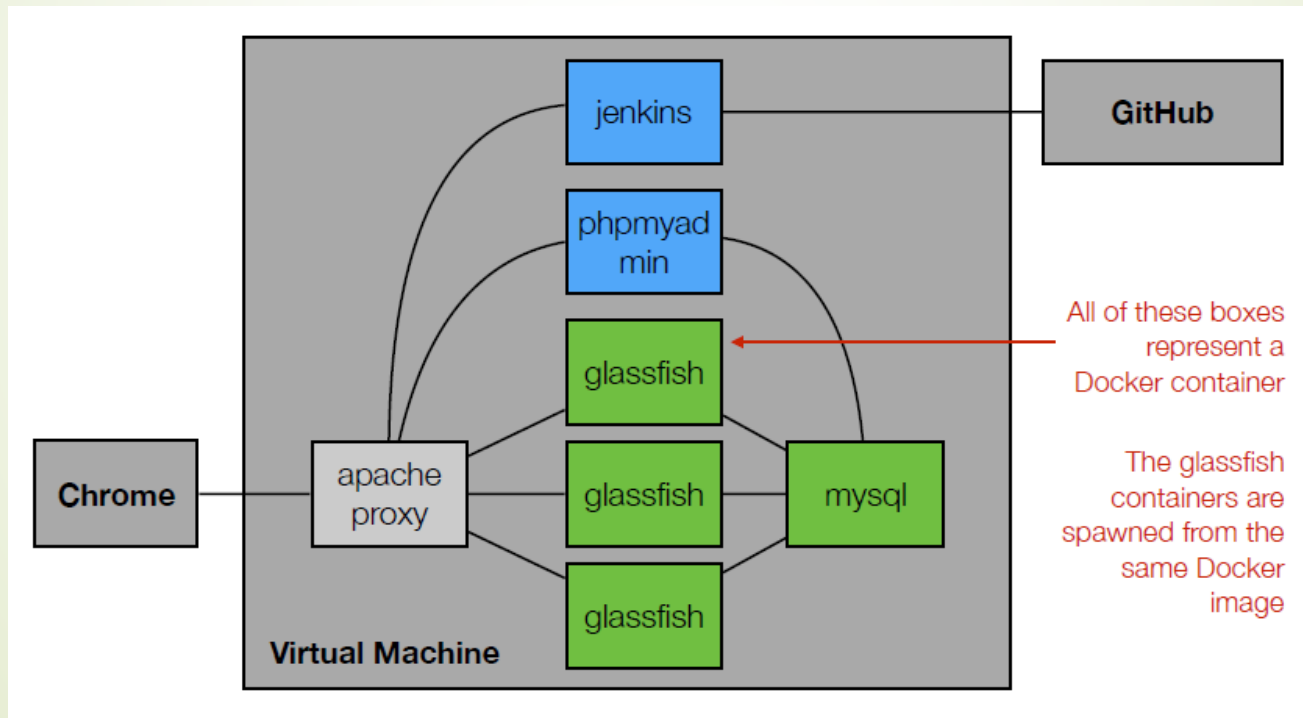
Introduction

- This presentation is about our lab on Vagrant and Docker
- The objective was to build an entire architecture of continuous integration and then deploy our first application on it
- Using two technologies:
 - Vagrant
 - Docker
- Repo : https://github.com/bbcnt/AMT_Vagrecker (private ATM)



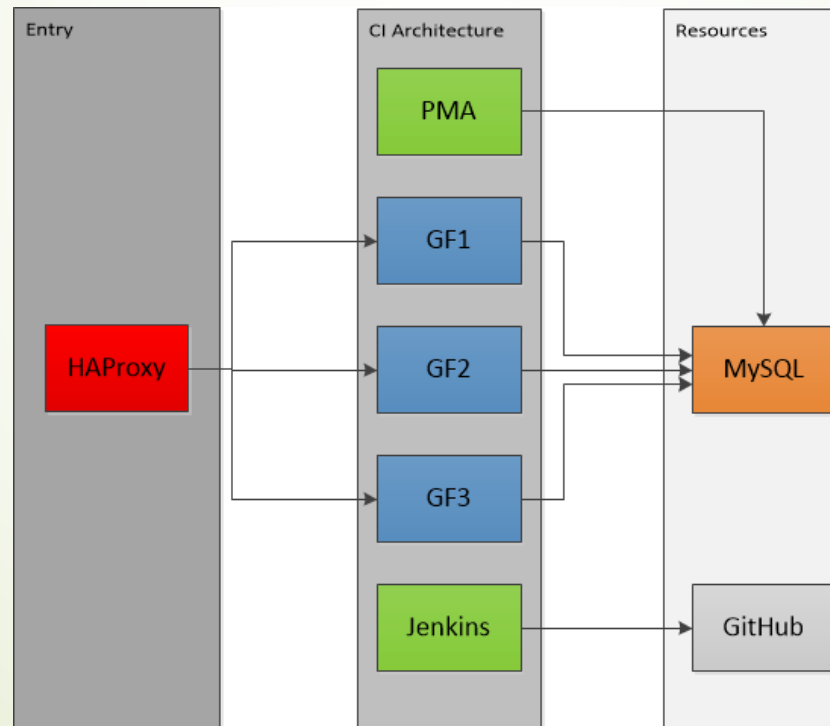
Architecture – The base


➤ This is the wanted result:



Architecture – Our version

- Let's take a look at the architecture that was built:





Architecture - Differences

- What's different ?
 - We replaced Apache by HAProxy
 - Why?
 - Because it was easier, and also, because it's a real Load Balancer
 - One problem remains though
 - We can't access the other services (Jenkins, PHPMyAdmin) with the same IP.
 - Example:
 - Jenkins: <http://localhost:9001>
 - PHPMyAdmin: <http://localhost:9004>
 - Web Application: <http://localhost:9002>
 - We make use of Docker's NAT system (mapping IP and Ports)



What we actually do

- Generate a .war file from the Maven project with Jenkins (the Maven project is pulled from GitHub)
- Deploy this project on a glassfish instance
- Link everything so the glassfish server can work with the MySQL database.
- Add a few tools, like PHPMyAdmin, to be able to work efficiently with MySQL.
- Add a few other instances of Glassfish (replicates of the first one).
- Put in some load balancing and we are done.

Sharing the volumes

- We use the sharing of volumes offered by Docker to make our stuff work.
- By sharing the autodeploy directory of gf1 with Jenkins and the other gf
 - (gf meaning glassfish, not girlfriend)

Containers Name	Shared volumes	Links	Volumes From
<i>mysql</i>	-	-	-
<i>phpmyadmin</i>	-	mysql	-
<i>gf1</i>	<i>/opt/glassfish/glassfish4/glassfish/domains/domain1/autodeploy/</i>	mysql	-
<i>gf2</i>	-	mysql	gf1
<i>gf3</i>	-	mysql	gf1
<i>jenkins</i>	-	gf1	gf1
<i>haproxy</i>	-	gf1 / gf2 / gf3	-

- We can make some magic work.

Sharing the volumes

- So, sharing this repo on gf1:
 - /opt/glassfish/glassfish4/glassfish/domains/domain1/autodeploy/
- Will create the same directory on those who make use of this volume.
- It will be very useful for deploying the application.

```
#Gf1 is sharing the autodeploy repertory, once a war file is put in this directory, it is automatically deployed by GF.
d.run "gf1", image: "heig/glassfish", args: "-p 4081:4848 -p 4082:8080 --link mysql:mysql -v /opt/glassfish/glassfish4/glassfish/domains/domain1/autodeploy/"
#And here, we use link the autoeplay repo of gf1, meaning, all the other GF servers will be updated once we update the first one.
d.run "gf2", image: "heig/glassfish", args: "-p 4083:4848 -p 4084:8080 --link mysql:mysql --volumes-from gf1"
d.run "gf3", image: "heig/glassfish", args: "-p 4085:4848 -p 4086:8080 --link mysql:mysql --volumes-from gf1"
```



Some Magic

- Glassfish offers a pretty cool thing, inside a domain, there is an “autodeploy” directory. Any war/jar/ear file inside will be deployed !
- So for Jenkins, it's easy, once it creates a war, it only need to copy it.
- Where?
 - Remember the shared volume?
 - By sharing the volumes from gf1 with the other gf instances and Jenkins, we only have to do this after the build on Jenkins:

```
cp $JENKINS_HOME/jobs/AMT_REST/workspace/target/AMT_REST-1.0-SNAPSHOT.war /opt/glassfish/glassfish4/glassfish/domains/domain1/autodeploy/AMT_REST.war
```
- And... that's all ;=)

Basic example

```
brito_000@BBC-LENOVO /F/VMs/AMT_Vagrant
$ vagrant up
Bringing machine 'default' up with 'virtualbox' provider...
==> default: Clearing any previously set forwarded ports...
==> default: Clearing any previously set network interfaces...
==> default: Preparing network interfaces based on configuration...
        default: Adapter 1: nat
        default: Adapter 2: hostonly
==> default: Forwarding ports...
        default: 7070 => 9001 (adapter 1)
        default: 9090 => 9002 (adapter 1)
        default: 3306 => 9003 (adapter 1)
        default: 5050 => 9004 (adapter 1)
        default: 4081 => 9005 (adapter 1)
        default: 4082 => 9006 (adapter 1)
        default: 4084 => 9007 (adapter 1)
        default: 4086 => 9008 (adapter 1)
        default: 4087 => 9009 (adapter 1)
        default: 22 => 2222 (adapter 1)
==> default: Booting VM...
==> default: Waiting for machine to boot. This may take a few minutes...
        default: SSH address: 127.0.0.1:2222
        default: SSH username: vagrant
        default: SSH auth method: private key
        default: Warning: Connection timeout. Retrying...
==> default: Machine booted and ready!
==> default: Checking for guest additions in VM...
==> default: Configuring and enabling network interfaces...
==> default: Mounting shared folders...
        default: /vagrant => F:/VMs/AMT_Vagrant
==> default: Machine already provisioned. Run 'vagrant provision' or use the '--provision'
==> default: to force provisioning. Provisioners marked to run always will still run.
```



Basic example

- We will now see how to make the AMT REST project work.
- First, we need to open a client in the Vagrant root directory
 - Where you'll find the Vagrantfile
- Start by typing:
 - `vagrant up`
 - `vagrant provision` (if not the first time)
- Then go to Jenkins : <http://localhost:9001>, click on the build button and wait.
- Access the app on http://localhost:9002/AMT_REST

Basic example

➤ Vagrant provision

```
==> default: ---> Using cache
==> default: ---> 5c86f9cd4a90
==> default: Step 2 : ADD haproxy.cfg /etc/haproxy/haproxy.cfg
==> default: ---> Using cache
==> default: ---> ac40235408f9
==> default: Step 3 : ADD start.bash /haproxy-start
==> default: ---> Using cache
==> default: ---> 9868e99b932f
==> default: Step 4 : VOLUME /haproxy-override
==> default: ---> Using cache
==> default: ---> a4ba7f0fed33
==> default: Step 5 : WORKDIR /etc/haproxy
==> default: ---> Using cache
==> default: ---> 6ae25177a61f
==> default: Step 6 : CMD bash /haproxy-start
==> default: ---> Using cache
==> default: ---> c2ba0b3d20ea
==> default: Step 7 : EXPOSE 80
==> default: ---> Using cache
==> default: ---> 093c41c71629
==> default: Step 8 : EXPOSE 443
==> default: ---> Using cache
==> default: ---> 9c0115c59b14
==> default: Successfully built 9c0115c59b14
==> default: Starting Docker containers...
==> default: -- Container: mysql
==> default: -- Container: phpmyadmin
==> default: -- Container: gf1
==> default: -- Container: gf2
==> default: -- Container: gf3
==> default: -- Container: jenkins
==> default: -- Container: haproxy
```

Basic example

The screenshot shows the Jenkins web interface running on localhost:9001. The browser's address bar displays 'http://localhost:9001'. The page features a sidebar on the left with navigation links: 'Nouveau Item', 'Utilisateurs', 'Historique des constructions', 'Administrer Jenkins', and 'Credentials'. The main content area displays a table of build jobs. The table has columns for 'S' (Status), 'M' (Icon), 'Nom du projet' (Project Name), 'Dernier succès' (Last Success), 'Dernier échec' (Last Failure), and 'Dernière durée' (Last Duration). A single job is listed with the name 'AMT_REST'. Below the table, there are links for 'Légende', 'RSS pour tout', 'RSS de tous les échecs', and 'RSS juste pour les dernières compilations'. At the bottom of the page, it says 'Page générée: 15 janv. 2015 07:10:00' and 'Jenkins ver. 1.595'. A red arrow points from the text 'Click Here' to the 'AMT_REST' job entry in the table.


S	M	Nom du projet ↓	Dernier succès	Dernier échec	Dernière durée
		AMT_REST	s. o.	s. o.	ND

Icône: [S](#) [M](#) [L](#)

[Légende](#) [RSS pour tout](#) [RSS de tous les échecs](#) [RSS juste pour les dernières compilations](#)

Page générée: 15 janv. 2015 07:10:00 [REST API](#) [Jenkins ver. 1.595](#)

Basic example



Vagrantfile

➤ Mapping des ports

```
#These are the port forwarding that are configured
config.vm.network "forwarded_port", guest: 7070, host: 9001 # Jenkins
config.vm.network "forwarded_port", guest: 9090, host: 9002 # Reverse proxy (entrypoint for users)
config.vm.network "forwarded_port", guest: 3306, host: 9003 # MySQL Server (pretty useless as is)
config.vm.network "forwarded_port", guest: 5050, host: 9004 # PHPMyAdmin (use localhost:9004/phpmyadmin)
config.vm.network "forwarded_port", guest: 4081, host: 9005 # GF1 config (4848)
config.vm.network "forwarded_port", guest: 4082, host: 9006 # GF1 web (8080)
config.vm.network "forwarded_port", guest: 4084, host: 9007 # GF2 web (8080)
config.vm.network "forwarded_port", guest: 4086, host: 9008 # GF3 web (8080)
config.vm.network "forwarded_port", guest: 4087, host: 9009 # HAProxy (stats on localhost:9009/haproxy login: admin, admin)
```


Vagrantfile

```
config.vm.provision "docker" do |d|

  #MySQL Server
  d.build_image "/vagrant/docker/mysq1", args: "-t heig/mysq1"
  d.run "mysq1", image: "heig/mysq1", args: "-p 3306:3306"

  #PHPMyAdmin (works with the MySQL server with the link)
  d.build_image "/vagrant/docker/phpmyadmin", args: "-t heig/phpmyadmin"
  d.run "phpmyadmin", image: "heig/phpmyadmin", args: "-p 5050:80 --link mysq1:mysq1"

  #Glassfish Servers

  d.build_image "/vagrant/docker/glassfish", args: "-t heig/glassfish"

  #Gf1 is sharing the autodeploy repertory, once a war file is put in this directory, it is automatically deployed by GF.
  d.run "gf1", image: "heig/glassfish", args: "-p 4081:4848 -p 4082:8080 --link mysq1:mysq1 -v /opt/glassfish/glassfish4/glassfish/domains/domain1/autodeploy/"
  #And here, we use link the autoeplody repo of gf1, meaning, all the other GF servers will be updated once we update the first one.
  d.run "gf2", image: "heig/glassfish", args: "-p 4083:4848 -p 4084:8080 --link mysq1:mysq1 --volumes-from gf1"
  d.run "gf3", image: "heig/glassfish", args: "-p 4085:4848 -p 4086:8080 --link mysq1:mysq1 --volumes-from gf1"


  #Jenkins Server

  #We create out .war file from the project, then we put it directly into gf1 autodeploy directory. We don't need to use
  #asadmin or anything else, it will be deployed automatically
  d.build_image "/vagrant/docker/jenkins", args: "-t heig/jenkins"
  d.run "jenkins", image: "heig/jenkins", args: "-p 7070:8080 --link gf1:gf1 --volumes-from gf1"

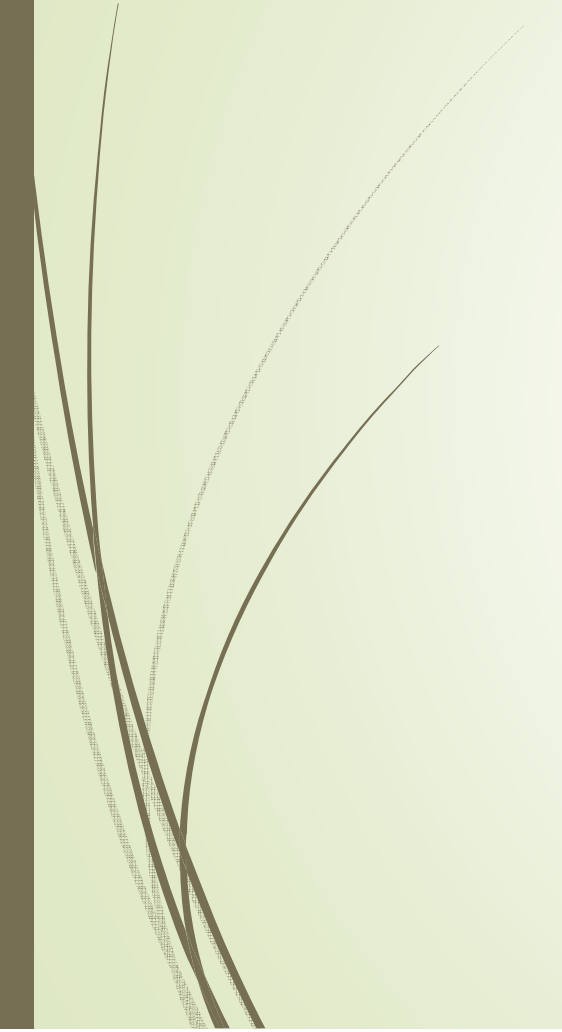
  #Reverse Proxy (technically, the user only needs to connect to localhost:9002 and he will access one of the 3 GF instances)

  d.build_image "/vagrant/docker/haproxy", args: "-t heig/haproxy"
  d.run "haproxy", image: "heig/haproxy", args: "-p 9090:80 -p 4087:8080 --link gf1:gf1 --link gf2:gf2 --link gf3:gf3"

end
```



Docker - MySQL



```
FROM dockerfile/ubuntu

RUN groupadd mysql
RUN useradd mysql -g mysql -u 4242

RUN apt-get update \
    && apt-get install -y mysql-server \
    && rm -rf /var/lib/mysql/mysql \
    && rm -rf /var/lib/apt/lists/* # 20140918

ADD start /start
ADD example_config.sql /tmp/example_config.sql
RUN chmod 755 /start

EXPOSE 3306

VOLUME ["/var/lib/mysql"]
VOLUME ["/run/mysqld"]

CMD ["/start"]
```

Docker - MySQL

- We create users that can remotely connect:

```
mysql -uroot -e "CREATE USER 'bruno'@'localhost' IDENTIFIED BY '123456';"  
mysql -uroot -e "CREATE USER 'bruno'@'%' IDENTIFIED BY '123456';"  
mysql -uroot -e "GRANT ALL PRIVILEGES ON * . * TO 'bruno'@'localhost';"  
mysql -uroot -e "GRANT ALL PRIVILEGES ON * . * TO 'root'@'%'";  
mysql -uroot -e "GRANT ALL PRIVILEGES ON * . * TO 'bruno'@'%'";
```

- And we add the PHPMyAdmin script in the DB

```
#Adding config file (some tables are needed for PHPMyAdmin)  
mysql -uroot < "/tmp/example_config.sql"
```

Docker - PHPMyAdmin

- Dockerfile too big, and not very useful. We install:
 - PHP
 - Apache
 - PHPMyAdmin
- Then interesting part is :

```
/* User used to manipulate with storage */  
$cfg['Servers'][$i]['controlhost'] = '';  
$cfg['Servers'][$i]['controluser'] = 'root';  
$cfg['Servers'][$i]['controlpass'] = '123456';
```

```
/*  
 * First server  
 */  
$i++;  
/* Authentication type */  
$cfg['Servers'][$i]['auth_type'] = 'cookie';  
/* Server parameters */  
$cfg['Servers'][$i]['host'] = 'mysql';  
$cfg['Servers'][$i]['connect_type'] = 'tcp';  
$cfg['Servers'][$i]['port'] = '3306';  
$cfg['Servers'][$i]['compress'] = false;  
/* Select mysql if your server does not have mysqli */  
$cfg['Servers'][$i]['extension'] = 'mysqli';  
$cfg['Servers'][$i]['AllowNoPassword'] = false;
```


Docker - Glassfish

- Again, too big to put here, but what we do:
 - We install Glassfish
 - We create the jdbc resource and pool linked to the MySQL container
 - We copy the mysql connector driver
 - And we start it.
- RUN \
 - `./asadmin start-domain && \`
 - `./asadmin --user admin --passwordfile pwdfile enable-secure-admin && \`
 - `./asadmin --user admin --passwordfile pwdfile create-jdbc-connection-pool --restype=javax.sql.XADataSource --datasourceclassname=com.mysql.jdbc.jdbc2.optional.MysqlXADataSource --property user=$DB_TECHNICAL_USER:password=$DB_TECHNICAL_USER_PASSWORD:serverName=mysql:portNumber=306:databaseName=$DB_NAME $JDBC_CONNECTION_POOL_NAME && \`
 - `./asadmin --user admin --passwordfile pwdfile create-jdbc-resource --connectionpoolid $JDBC_CONNECTION_POOL_NAME $JDBC_JNDI_NAME && \`
 - `./asadmin stop-domain`



Docker - Jenkins

- What we do:
 - We install Java and Maven (to be able to create Maven projects)
 - We get Jenkins to work (install and setting environment)
 - We add the config files where we want and need them (we'll see those in a moment)
 - We also add the needed plugins (mostly for github support)
 - And that's all
- We also need to save a few config files (xml).



Docker - HAProxy

```
FROM dockerfile/ubuntu
# Install Haproxy.
RUN \
sed -i 's/^# \(\.*-backports\s\)/\1/g' /etc/apt/sources.list && \
apt-get update && \
apt-get install -y haproxy && \
sed -i 's/^ENABLED=.*ENABLED=1/' /etc/default/haproxy && \
rm -rf /var/lib/apt/lists/*
# Add files.
ADD haproxy.cfg /etc/haproxy/haproxy.cfg
ADD start.bash /haproxy-start
# Define mountable directories.
VOLUME ["/haproxy-override"]
# Define working directory.
WORKDIR /etc/haproxy
# Define default command.
CMD ["bash", "/haproxy-start"]
# Expose ports.
EXPOSE 80
EXPOSE 443
```

Docker - HAProxy

- And setting HAProxy
- Remember this:

```
root@9b270c07454b:/etc/haproxy# cat /etc/hosts
172.17.0.8      9b270c07454b
127.0.0.1      localhost
::1            localhost ip6-localhost ip6-loopback
fe00::0        ip6-localnet
ff00::0        ip6-mcastprefix
ff02::1        ip6-allnodes
ff02::2        ip6-allrouters
172.17.0.4      gf1
172.17.0.5      gf2
172.17.0.6      gf3
```

```
...
frontend http_stats
  bind :8080
  stats enable
  stats uri      /haproxy
  stats realm    Haproxy\ Statistics
  stats auth     admin:admin

frontend http_proxy
  bind :80
  default_backend gf_servers

backend gf_servers
  server gf1 gf1:8080/AMT_REST check
  server gf2 gf2:8080/AMT_REST check
  server gf3 gf3:8080/AMT_REST check
```

- We use the address given by the Vagrant VM.

Docker - HAProxy

- And that's all, we can now work with the API.
- Other cool thing, stats with HAProxy:

http://localhost:9009/haproxy

HAProxy version 1.4.24, released 2013/06/17

Statistics Report for pid 7

> General process information

pid = 7 (process #1, nbproc = 1)
uptime = 0d 0h10m57s
system limits: memmax = unlimited; ulimit-n = 4015
maxsock = 4015; maxconn = 2000; maxpipes = 0
current conns = 1; current pipes = 0/0
Running tasks: 1/4

active UP, active UP, going down, active DOWN, going up, active or backup DOWN, active or backup DOWN for maintenance (MAINT), backup UP, backup UP, going down, backup DOWN, going up, not checked

Display option:
• [Hide DOWN servers](#)
• [Refresh now](#)
• [CSV export](#)

External resources:
• [Primary site](#)
• [Updates \(v1.4\)](#)
• [Online manual](#)

Note: UP with load-balancing disabled is reported as "NOLB".

http_stats

	Queue			Session rate			Sessions				Total	LbTot	Bytes		Denied		Errors		Warnings		Server									
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	In			Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Downtme	Thrtle	
Frontend	1	4	-	1	1	2 000	22			18 013	249 637	0	0	0						OPEN										

http_proxy

	Queue			Session rate			Sessions				Total	LbTot	Bytes		Denied		Errors		Warnings		Server									
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	In			Out	Req	Resp	Req	Conn	Resp	Retr	Redis	Status	LastChk	Wght	Act	Bck	Chk	Dwn	Downtme	Thrtle	
Frontend	0	1	-	0	4	2 000	33			6 603	12 628	0	0	0						OPEN										

gf_servers

	Queue			Session rate			Sessions				Total	LbTot	Bytes		Denied		Errors		Warnings		Status	LastChk	Server									
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	In			Out	Req	Resp	Req	Conn	Resp	Retr	Redis			Wght	Act	Bck	Chk	Dwn	Downtme	Thrtle			
gf1	0	0	-	0	1		0	1	-	11	11	2 169	4 753	0	0	0	0	0	0	10m35s UP	L4OK in 0ms	1	Y	-	0	1	22s	-				
gf2	0	0	-	0	1		0	1	-	11	11	2 215	4 006	0	0	0	0	0	0	10m34s UP	L4OK in 0ms	1	Y	-	0	1	22s	-				
gf3	0	0	-	0	1		0	2	-	11	11	2 219	3 869	0	0	0	0	0	0	10m32s UP	L4OK in 0ms	1	Y	-	0	1	24s	-				
Backend	0	0		0	1		0	4	0	33	33	6 603	12 628	0	0	0	0	0	0	10m35s UP		3	3	0		1	21s					



A few tricks

- A few cool things we noted.
 - Connecting to a docker:
 - `sudo docker exec -i -t container_name bash`
 - Download of big files (max of GitHub is 50 Mo), here jdk 8u25
 - `curl -v -j -k -L -H "Cookie: oraclelicense=accept-securebackup-cookie" http://download.oracle.com/otn-pub/java/jdk/8u25-b17/jdk-8u25-linux-x64.rpm > jdk-8u25-linux-x64.rpm`
 - `wget --no-check-certificate --no-cookies --header "Cookie: oraclelicense=accept-securebackup-cookie" http://download.oracle.com/otn-pub/java/jdk/8u25-b17/jdk-8u25-linux-x64.rpm`
 - Otherwise it won't work in a script ;)

Conclusion

- We can't show you everything, but as you can see, not so hard.
- Took a lot of time to actually configure the dockers, more than setting them
- We also tried to make this as light as possible
 - By downloading files inside the dockers and not keeping them in the host
- Vagrant and Docker were supposed to be concurrent
 - In the end, they work pretty well together!

